

What are the most important variables
that determine a person's income?

Fabrizio Stella

Abstract

The purpose of this work is to understand **what are the individual characteristics of people who better than others explain the differences in wages among people.**

The **dataset used** in this study **come from** the “**UCI Machine Learning Repository**” and is based on the 1994 census. It's composed by variables like age, education, occupation, sex, race, birth country, salary and others.

To study this relation a **decision tree classification model** has been created and performed by using **Scikit-learn** library.

The results seem to suggest that a **relation between individual characteristics of people and salary exists**, and the model built seems to predict correctly more than 80% of the data entered. The variables with the most informative capacity are: **Relationship, Education, Occupation and Age**. Other variables like race, sex, native country are not significant.

Motivation

What is driving my interest in the past months are the variables – personal or not – that can explain the future income of people. Many of them, of course, are not under our control: age or sex could explain income differences across people, for example, but aren't under our control. Others, instead, at least partially, are under our control: Education for example.

Being able to determine the contribution of these variables to the future income of people is very important as we can better drive the behavior of the youngest to the right direction, telling them how much important is, if it is, education for their future or, at worst, we can, in any case, extract relevant policy implication.

Dataset(s)

The dataset used in this data science project is called “Adult” and come from the well-known data science repository “UCI Machine Learning Repository”. It was created in 1996 in the USA and it’s based on the 1994 Census Database.

It’s a collection of individual characteristics of people, including age, education, occupation, race, salary. The number of instances at the beginning of the analysis was 32560. The number of categorical variables is 15: among them, salary is our Target Variable.

The underlying idea behind the dataset is that a relation between Salary and all the other variables can be established.

Data Preparation and Cleaning

As said before, the dataset I used for my analysis is “Adult”, from the “UCI Machine Learning Repository”. At the beginning of the analysis, the dataset was composed by 15 attributes (including the target one) and 32560 instances.

In the figure 1 is shown the list of the attributes in the initial dataset.

Figure 1

Variable Name	Kind	Description
age	continuous	17<age<90
workclass	categorical	8 different unique instances
fnlwgt	continuous	
education	categorical	16 different unique instances
education-num	continuous	16 different numbers
marital-status	categorical	7 different unique instances
occupation	categorical	14 different unique instances
relationship	categorical	6 different unique instances
race	categorical	5 different unique instances
sex	categorical	2 different unique instances
capital-gain	continuous	
capital-loss	continuous	
hours-per-week	continuous	1<hours per week<99
native-country	categorical	40 different unique instances
Salary	categorical	2 different unique instances > 50K\$ < 50K\$

Main actions taken to clean the dataset:


1. Drop not relevant columns: **the columns “fnlwgt”, education-num, capital-gain and capital-loss have been removed** from the dataset. The variable “fnlwgt” is considered not relevant for the analysis purposes. In capital-gain and capital-loss, instead, too many missing values made impossible their use for further analysis.
2. All rows in which were at least 1 **NaN values have been removed** from the dataset. This implies a reduction of 8% of the number of instances (from 32560 to 30117).
3. Attribute manipulation: for not removed attributes a further cleaning activity has been performed: **continuous variable have been transformed into categorical ones** and the **number of distinct observations of categorical variables has been reduced**. In the next slides some example of this activity will be provided.

Target Variable


Data Preparation and Cleaning. Focus on attributes manipulation:

Reduction of the distinct values assumed by the 'Age' Attribute.

Before
Changes




After
Changes




Data Preparation and Cleaning: Focus on attributes manipulation:

Reduction of the distinct values assumed by the 'Working Hours per Week' Attribute.

Before
Changes



After
Changes



Research Question(s)

In this project I'm trying to understand what are the main individual characteristics associated to high salaries. Variables like age, education, race, sex, could somehow affect a person's salary. I would like to understand if this is true and, if it is, which of these variables are the most important.

Through a classification task I'll verify whether the behavior of some of the variables mentioned above could be associated with the behavior of the remuneration that a person receive from the market.


In other, simpler word, what are the most important individual characteristics that explain the differences between individuals in the probability of getting a better wage? Is it true that educated people receive higher salaries? And, is it true that women, on average, receive lower salaries than men?

Methods

For this analysis I used **Scikit learn** to perform a **decision tree data classification**, as I know ex-ante the values of the target variable and many attributes are categorical (including the Target one).

To perform this decision tree the following actions have been taken:

1. I divided the **dataset into two**:

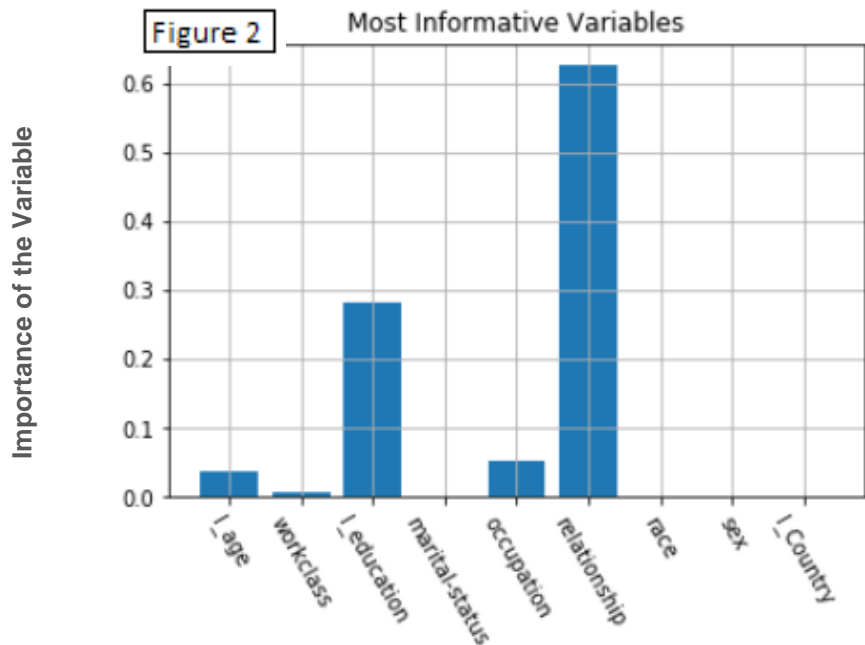


2. I applied the decision tree data classification algorithm to the **Training set** with the following parameters:

- ✓ **No maximum depth** has been imposed
- ✓ **Maximum leaf node** == 20

Findings: Results from the Training Phase (1)

After having applied the previously traced model during the training phase, we can determine how much good our attribute are and, in particular, wich are the most informative among them, i.e. the attributes that better than other help us in understand the behavior of our target variable: salary.



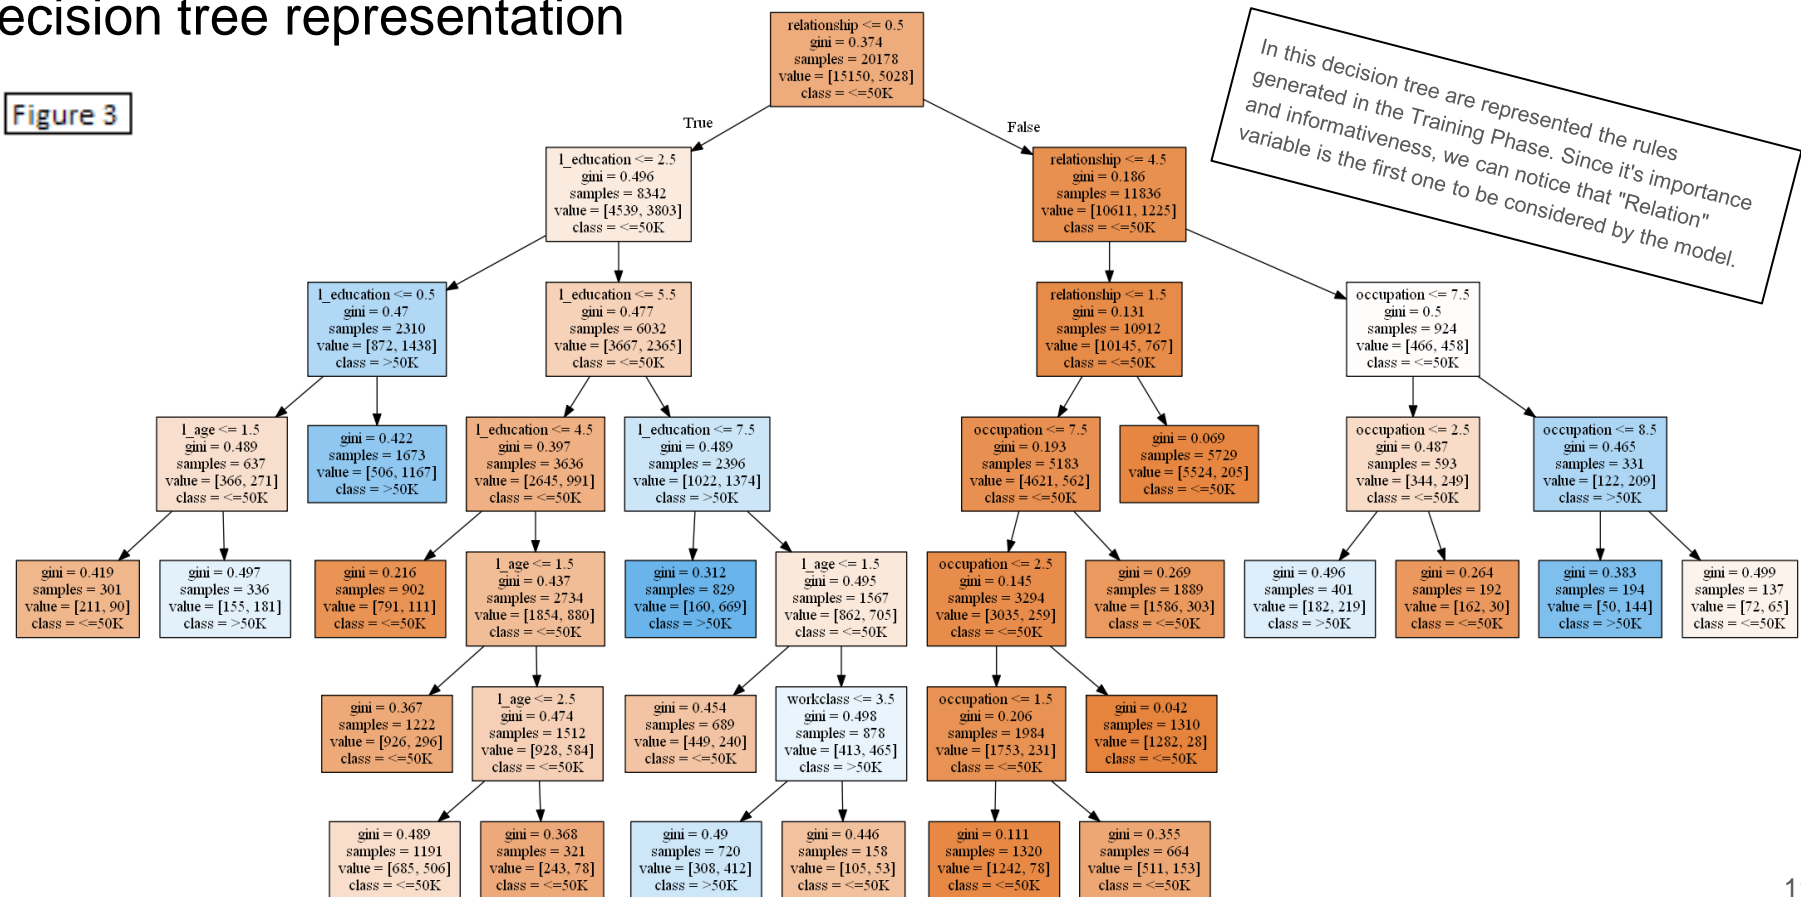
❖ In the **Figure 2** are displayed the **variables** that compose our **dataset**. The information we can get is the level of **importance** of each of them in terms of **informativeness**.

- ✓ Surprisingly, “**relationship**” is the most informative variable. It includes observations like “Unmarried”, “Husband”, “Wife”, “Own-Child”, “not in Family”, “Other Relative”.
- ✓ The other most informative variables are **Education**, **Occupation**, **age**, **workclass**.
- ✓ No importance is associated to the other variables: Marital Status, Race, Sex, Country. They seems no to be able to explain income variation across people.

Findings: Results from the Training Phase (2):

Decision tree representation

Figure 3



Findings: Results from the Test Phase

- ✓ The last point remained is: how good the built model is. Is it able to predict the target variable?
- ✓ To answer this question we can use the test dataset previously generated and try to understand how many mistakes we get by the application of the training phase model.
- ✓ One measure we can use the accuracy of our model is the '**Accuracy Score**':
 - It is equal to 1 if the model is perfect;
 - It is equal to 0 if the model is totally mistaken and no variable is classified correctly;
- ✓ An accuracy score of .75 could be considered good.

Figure 4

Measure of accuracy

```
In [775]: accuracy_score(y_true = y_test,  
                        y_pred = predictions)
```

```
Out[775]: 0.820303853506389
```

In the Figure 4 is displayed the accuracy score: it is .82. It's very good: It's means that more than 82% of the instances that make up the Test set are classified correctly.

Limitations

This project work and its results have **some relevant limits**:

- 1) It's based on a census done in a town in the USA: it implies that probably it could not be applicable to other different places or Countries.
- 2) I'm not creating a model of causation: We can say that, based on this model, higher education is associated with a higher probability to get a salary >50K; but this is not a causation: I cannot say that higher education cause higher salaries. I think that is true, of course; but this is not the instrument that prove this statement. For example, richest families can finance better education to their children. In this case would not be education to cause good salaries but initial conditions like parent's money.
- 3) Since the target variable is categorical, the dataset is divided into two, based on the values of the target variable which are either >50k\$ or <50k\$. This, near the threshold value, generates identification problems for individuals: a person with an income of 49.5k\$ should not be too much different from another with an income of 51.5k\$. But the model, as it is built, divides them totally. Unfortunately, not having the exact data of income, using a classification model is the only method to derive some information.

Conclusions

- ✓ Using Scikit-learn, I created a decision tree classification model that **predict the relation between individual characteristics and Salary**.
- ✓ **The level of accuracy of this model is roughly 82%:** it means that, knowing ex-ante some individual characteristics, we can predict whether a person will have an income higher than 50K\$ or not.
- ✓ Among the others, the most important variables that explain differences in salaries are **Relationship, Education, Occupation and Age**. Remain variables like race, sex, native country seem to be no useful in this.

Acknowledgements

Data used come from the repository called UCI Machine Learning repository, suggested by the professor Leo during the week 9 of the course of Python for Data Science.

All the analysis done in this final project come from my effort.

References

<https://www.python.org/doc/>

<https://pandas.pydata.org/docs/>

<https://stackoverflow.com/>

<https://medium.com/>

<https://scikit-learn.org/>

<https://numpy.org/>

<https://www.graphviz.org/>

Learning Python, 5th Edition, Mark Lutz

In [1]:

```
import pandas as pd
import numpy as np
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
#pd.set_option('display.max_rows', None)
pd.options.mode.chained_assignment = None
%matplotlib inline
import matplotlib.pyplot as plt
from sklearn import preprocessing
from matplotlib import pyplot
from numpy import loadtxt
from xgboost import XGBClassifier
from xgboost import plot_importance
import os
import graphviz
import pydotplus
from sklearn import tree
```

In [2]:

```
os.environ['PATH'] += os.pathsep + r'C:\Program Files\graphviz-2.38\release\bin'
```

In [3]:

```
#conda install -c anaconda py-xgboost
```

In [4]:

```
data = pd.read_csv('./AdultDataset/adult.data')
```

In [5]:

```
data.columns
```

Out[5]:

```
Index(['39', ' State-gov', ' 77516', ' Bachelors', ' 13', ' Never-married',
      ' Adm-clerical', ' Not-in-family', ' White', ' Male', ' 2174', ' 0',
      ' 40', ' United-States', ' <=50K'],
      dtype='object')
```

Description of the main variables in the dataset

1) age: continuous. 2) workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked. 3) fnlwgt: continuous. 4) education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool. 5) education-num: continuous. 6) marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse. 7) occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces. 8) relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried. 9) race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black. 10) sex: Female, Male. 11) capital-gain: continuous. 12) capital-loss: continuous. 13) hours-per-week: continuous. 14) native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad&Tobago, Peru, Hong, Holand-Netherlands.

Below, I insert the columns names

In [6]:

```
column_names = ['age']
```

```
column_names= [ 'age',
                 'workclass',
                 'fnlwgt',
                 'education',
                 'education-num',
                 'marital-status',
                 'occupation',
                 'relationship',
                 'race',
                 'sex',
                 'capital-gain',
                 'capital-loss',
                 'hours-per-week',
                 'native-country',
                 'Salary'
               ]

df1 = pd.DataFrame(data.values, columns = column_names )
```

In [7]:

```
df1.head()
```

Out[7]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss	hours-per-week	native-country
0	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	13	United-States
1	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0	0	40	United-States
2	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	0	0	40	United-States
3	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	0	40	Cuba
4	37	Private	284582	Masters	14	Married-civ-spouse	Exec-managerial	Wife	White	Female	0	0	40	United-States

I remove education num, as It's gives me not additional information compared to Education column

In [8]:

```
del(df1['education-num'])
```

In [9]:

```
df1.columns
```

Out[9]:

```
Index(['age', 'workclass', 'fnlwgt', 'education', 'marital-status',
       'occupation', 'relationship', 'race', 'sex', 'capital-gain',
       'capital-loss', 'hours-per-week', 'native-country', 'Salary'],
      dtype='object')
```

In [10]:

```
df1.shape
```

Out[10]:

```
(32560, 14)
```

In [11]:

```
df1.isnull().any()
```

Out[11]:

```
age                False
workclass           False
fnlwgt             False
education           False
marital-status      False
occupation          False
relationship        False
race               False
sex                False
capital-gain        False
capital-loss        False
hours-per-week      False
native-country      False
Salary             False
dtype: bool
```

In [12]:

```
toBeStrip=['workclass','education','marital-status','occupation',
            'relationship','race','sex','native-country','Salary']
for elem in toBeStrip:
    df1[elem]=df1[elem].str.strip()
```

Studying Single attributes

In [13]:

```
df1.columns
```

Out[13]:

```
Index(['age', 'workclass', 'fnlwgt', 'education', 'marital-status',
       'occupation', 'relationship', 'race', 'sex', 'capital-gain',
       'capital-loss', 'hours-per-week', 'native-country', 'Salary'],
      dtype='object')
```

Age

In [14]:

```
print("The Minimum Age is: ", df1['age'].min(),";",
      "The maximum age is: ", df1['age'].max(),";",
      "The average age is: ", df1['age'].mean(),";",
      "The mode age is: ", df1['age'].mode())
```

```
The Minimum Age is: 17 ; The maximum age is: 90 ; The average age is: 38.581633906633904 ; The
mode age is: 0 36
dtype: object
```

In [15]:

```
df1[df1['age']=='?'].any()
```

Out[15]:

```
age                False
workclass           False
fnlwgt             False
education           False
marital-status      False
occupation          False
relationship        False
```

```
race                False
sex                 False
capital-gain        False
capital-loss        False
hours-per-week      False
native-country      False
Salary              False
dtype: bool
```

In [16]:

```
from collections import Counter
l=[]
for elem in df1['age']:
    l.append(elem)
#Counter(l)
```

- I can create 4-5 variable for the attribute Age: <20; >20<40; >40<60; >60<80; >80

workclass

In [17]:

```
df1['workclass'].value_counts()
```

Out[17]:

```
Private                22696
Self-emp-not-inc       2541
Local-gov              2093
?                      1836
State-gov              1297
Self-emp-inc           1116
Federal-gov            960
Without-pay            14
Never-worked           7
Name: workclass, dtype: int64
```

In [18]:

```
df1[df1['workclass']=='?'].any()
```

Out[18]:

```
age                True
workclass           True
fnlwgt             True
education           True
marital-status      True
occupation          True
relationship        True
race                True
sex                 True
capital-gain        True
capital-loss        True
hours-per-week      True
native-country      True
Salary              True
dtype: bool
```

- What can I do with "Without pay" and "Never Worked?" Moreover I have to remove rows with "?"

Education

In [19]:

```
df1['education'].value_counts()
```

Out[19]:

```
HS-grad      10501
Some-college  7291
Bachelors    5354
Masters      1723
Assoc-voc    1382
11th         1175
Assoc-acdm   1067
10th         933
7th-8th      646
Prof-school  576
9th          514
12th         433
Doctorate    413
5th-6th      333
1st-4th      168
Preschool    51
Name: education, dtype: int64
```

- It has to be reduced the number of levels of instruction classes

In [20]:

```
df1[df1['education']=='?'].any()
```

Out[20]:

```
age          False
workclass     False
fnlwgt       False
education     False
marital-status False
occupation    False
relationship  False
race         False
sex          False
capital-gain  False
capital-loss  False
hours-per-week False
native-country False
Salary       False
dtype: bool
```

Marital Status

In [21]:

```
df1['marital-status'].value_counts()
```

Out[21]:

```
Married-civ-spouse    14976
Never-married         10682
Divorced              4443
Separated             1025
Widowed              993
Married-spouse-absent  418
Married-AF-spouse      23
Name: marital-status, dtype: int64
```

- Should I eliminate "married spouse absent" and "married af spouse?"

In [22]:

```
df1[df1['marital-status']=='?'].any()
```

Out[22]:

```
Out[22]:
```

```
age                False
workclass           False
fnlwgt             False
education           False
marital-status     False
occupation          False
relationship        False
race               False
sex                False
capital-gain        False
capital-loss        False
hours-per-week      False
native-country      False
Salary             False
dtype: bool
```

Occupation

```
In [23]:
```

```
df1['occupation'].value_counts()
```

```
Out[23]:
```

```
Prof-specialty      4140
Craft-repair        4099
Exec-managerial     4066
Adm-clerical        3769
Sales               3650
Other-service       3295
Machine-op-inspct   2002
?                  1843
Transport-moving    1597
Handlers-cleaners   1370
Farming-fishing     994
Tech-support        928
Protective-serv     649
Priv-house-serv     149
Armed-Forces         9
Name: occupation, dtype: int64
```

```
In [24]:
```

```
df1['occupation'].value_counts().shape[0]
```

```
Out[24]:
```

```
15
```

- I can remove "Armed-Forces" and "?"

```
In [25]:
```

```
df1[df1['occupation']!='?'].any()
```

```
Out[25]:
```

```
age                True
workclass           True
fnlwgt             True
education           True
marital-status     True
occupation          True
relationship        True
race               True
sex                True
capital-gain        True
capital-loss        True
hours-per-week      True
```

```
native-country    True
Salary            True
dtype: bool
```

relationship

In [26]:

```
df1['relationship'].value_counts()
```

Out[26]:

```
Husband          13193
Not-in-family     8304
Own-child         5068
Unmarried         3446
Wife              1568
Other-relative    981
Name: relationship, dtype: int64
```

In [27]:

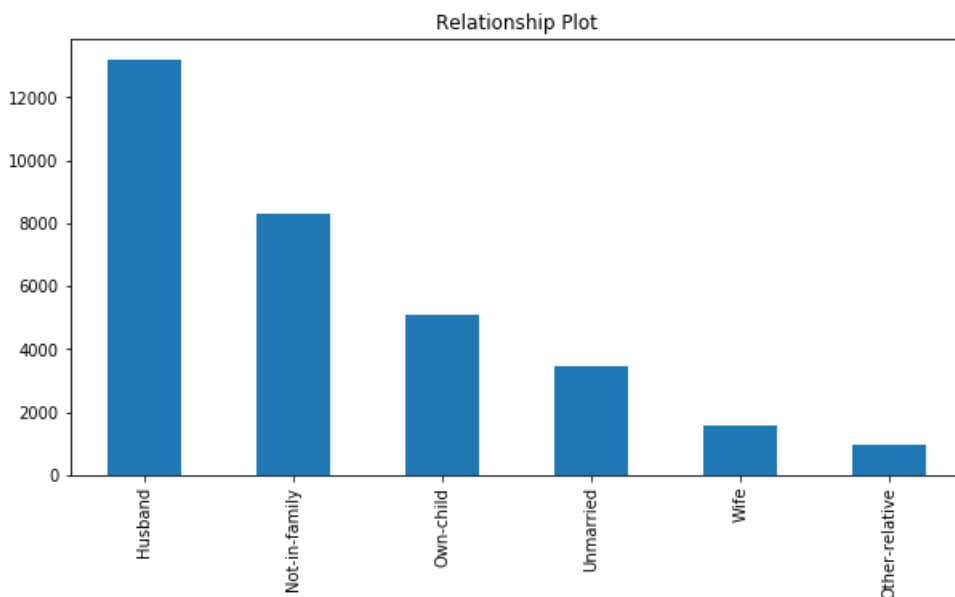
```
df1[df1['relationship']=='?'].any()
```

Out[27]:

```
age                False
workclass           False
fnlwgt             False
education           False
marital-status      False
occupation          False
relationship        False
race               False
sex                False
capital-gain        False
capital-loss        False
hours-per-week      False
native-country      False
Salary             False
dtype: bool
```

In [28]:

```
RelationshipPlot=df1['relationship'].sort_values().value_counts().plot.bar(figsize=(10,5),x='Value', title="Relationship Plot")
```



race

In [29]:

```
df1['race'].value_counts()
```

Out[29]:

```
White          27815
Black           3124
Asian-Pac-Islander  1039
Amer-Indian-Eskimo    311
Other            271
Name: race, dtype: int64
```

In [30]:

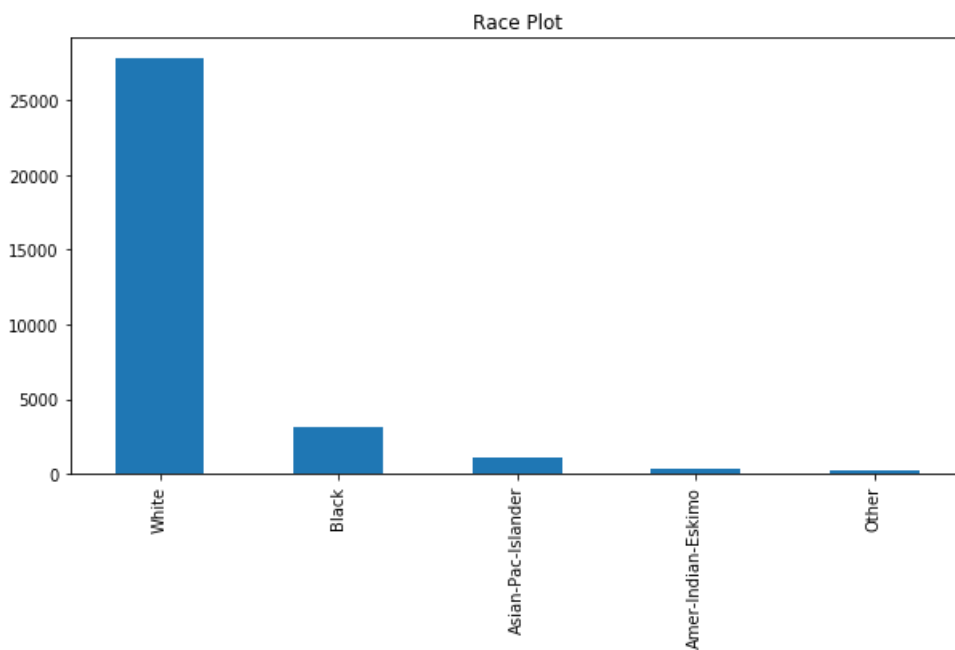
```
df1[df1['race']=='?'].any()
```

Out[30]:

```
age           False
workclass     False
fnlwgt        False
education     False
marital-status False
occupation   False
relationship  False
race          False
sex           False
capital-gain  False
capital-loss  False
hours-per-week False
native-country False
Salary        False
dtype: bool
```

In [31]:

```
RacePlot=df1['race'].sort_values().value_counts().plot.bar(figsize=(10,5),x='Value', title="Race Plot")
```



Sex

In [32]:


```
df1['sex'].value_counts()
```

Out[32]:

```
Male      21789
Female    10771
Name: sex, dtype: int64
```

In [33]:

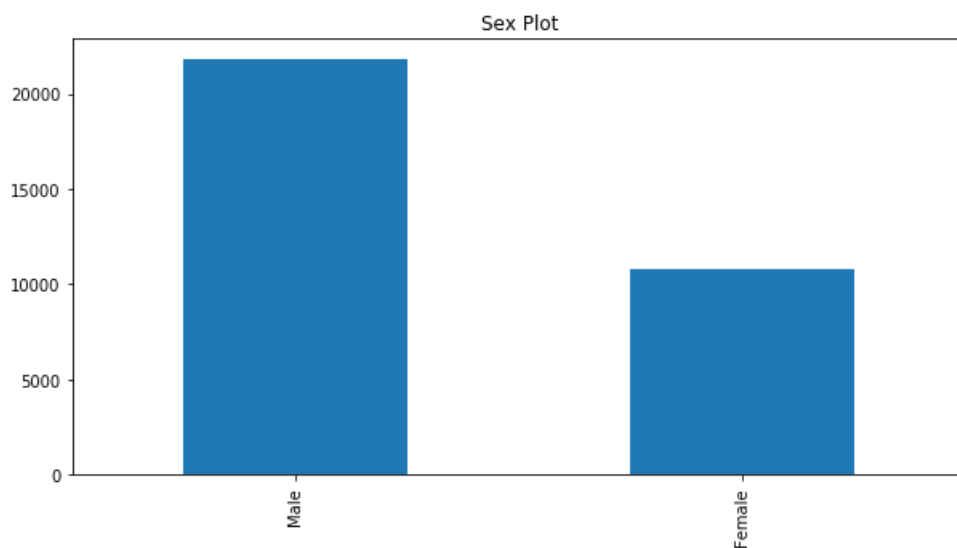
```
df1[df1['sex']=='?'].any()
```

Out[33]:

```
age          False
workclass     False
fnlwgt       False
education     False
marital-status False
occupation    False
relationship  False
race         False
sex          False
capital-gain  False
capital-loss  False
hours-per-week False
native-country False
Salary       False
dtype: bool
```

In [34]:

```
SexPlot=df1['sex'].sort_values().value_counts().plot.bar(figsize=(10,5),x='Value', title="Sex Plot"
)
```



capital-gain & Capital loss

In [35]:

```
df1['capital-gain'].value_counts()[:5]
```

Out[35]:

```
0      29849
15024   347
7688    284
7298    246
99999   159
Name: capital-gain, dtype: int64
```

In [36]:

```
df1['capital-loss'].value_counts()[ :5]
```

Out[36]:

```
0          31041
1902         202
1977         168
1887         159
1485          51
Name: capital-loss, dtype: int64
```

In [37]:

```
df1[df1['capital-loss']==0].shape[0]
```

Out[37]:

```
31041
```

In [38]:

```
df1[df1['capital-gain']==0].shape[0]
```

Out[38]:

```
29849
```

- Probably Both to be removed

hours-per-week

In [39]:

```
df1['hours-per-week'].value_counts()[ :5]
```

Out[39]:

```
40      15216
50       2819
45       1824
60       1475
35       1297
Name: hours-per-week, dtype: int64
```

In [40]:

```
df1['hours-per-week'].min()
```

Out[40]:

```
1
```

In [41]:

```
df1['hours-per-week'].max()
```

Out[41]:

```
99
```

In [42]:

```
df1['hours-per-week'].mean()
```

```
df1['hours-per-week'].mean()
```

Out[42]:

40.437469287469284

In [43]:

```
df1['hours-per-week'].mode()
```

Out[43]:

0 40
dtype: object

- This variable has to be arranged

In [44]:

```
df1[df1['hours-per-week']=='?'].any()
```

Out[44]:

age	False
workclass	False
fnlwgt	False
education	False
marital-status	False
occupation	False
relationship	False
race	False
sex	False
capital-gain	False
capital-loss	False
hours-per-week	False
native-country	False
Salary	False

dtype: bool

Native-Country

In [45]:

```
df1['native-country'].value_counts()[:10]
```

Out[45]:

United-States	29169
Mexico	643
?	583
Philippines	198
Germany	137
Canada	121
Puerto-Rico	114
El-Salvador	106
India	100
Cuba	95

Name: native-country, dtype: int64

- Devo creare variabile others

In [46]:

```
df1[df1['native-country']=='?'].any()
```

Out[46]:

age	True
workclass	True
fnlwgt	True
education	True
marital-status	True
occupation	True
relationship	True
race	True
sex	True
capital-gain	True
capital-loss	True
hours-per-week	True
native-country	True
Salary	True

```
workclass      True
fnlwgt         True
education      True
marital-status True
occupation     True
relationship   True
race           True
sex            True
capital-gain   True
capital-loss   True
hours-per-week True
native-country True
Salary         True
dtype: bool
```

Salary

In [47]:

```
df1['Salary'].value_counts()
```

Out[47]:

```
<=50K    24719
>50K      7841
Name: Salary, dtype: int64
```

In [48]:


```
df1[df1['Salary']=='?'].any()
```

Out[48]:

```
age           False
workclass     False
fnlwgt        False
education     False
marital-status False
occupation    False
relationship   False
race          False
sex           False
capital-gain   False
capital-loss   False
hours-per-week False
native-country False
Salary        False
dtype: bool
```

In [49]:

```
SalaryPlot=df1['Salary'].sort_values().value_counts().plot.bar(figsize=(10,5),x='Value', title="Salary Plot")
```



<=50K

>50K

To do

1. Convert Unknown to "?" and eliminate it

In [50]:

```
beforeRows = df1.shape[0]
print(beforeRows)
```

32560

I eliminate all the rows in which there is at least a missing value in a certain attribute

Thanks to the prior analysis i know that missing values are in these three instances: 1) workclass; 2) occupation; 3) native-country

First Filter

In [51]:

```
df1=df1[df1['workclass']!='?']
```

In [52]:

```
df1.shape
```

Out[52]:

(30724, 14)

Second Filter

In [53]:

```
df1[df1['occupation']=='?'].any()
```

Out[53]:

age	True
workclass	True
fnlwgt	True
education	True
marital-status	True
occupation	True
relationship	True
race	True
sex	True
capital-gain	False
capital-loss	False
hours-per-week	True
native-country	True

```
Salary          True
dtype: bool
```

```
In [54]:
```

```
df1=df1[df1['occupation']!='?']
```

```
In [55]:
```

```
df1.shape
```

```
Out[55]:
```

```
(30717, 14)
```

Third Filter

```
In [56]:
```

```
df1=df1[df1['native-country']!='?']
```

```
In [57]:
```

```
df1.shape
```

```
Out[57]:
```

```
(30161, 14)
```

Here the remaining columns

```
In [58]:
```

```
afterRows = df1.shape[0]
print(afterRows)
```

```
30161
```

How many rows dropped due to cleaning?

```
In [59]:
```

```
beforeRows - afterRows
```

```
Out[59]:
```

```
2399
```

```
In [60]:
```

```
df1[:1]
```

```
Out[60]:
```

	age	workclass	fnlwgt	education	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss	hours-per-week	native-country	Salary
0	50	Self-emp-not-inc	83311	Bachelors	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	13	United-States	<=50K

Generic code to make discrete all the cotinuous variables

Age

In [61]:

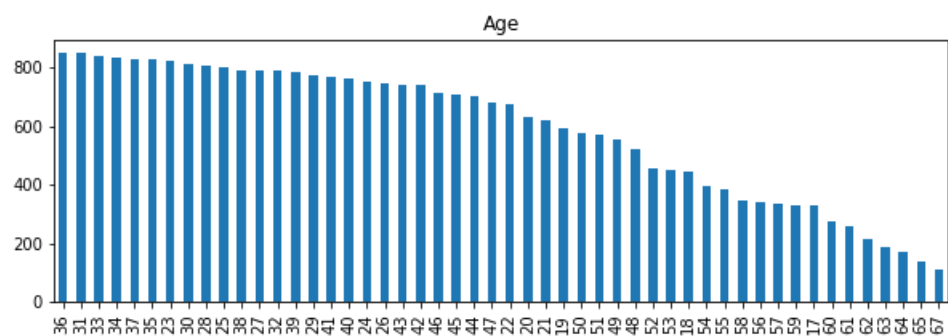
```
l_age=[]
for elem in df1['age']:
    if elem <20:
        l_age.append('<20')
    elif elem >20 and elem <40:
        l_age.append('>20 & <40')
    elif elem >40 and elem <60:
        l_age.append('>40 & <60')
    elif elem >60 and elem <80:
        l_age.append('>60 & <80')
    else:
        l_age.append('>80')
```

In [62]:

```
df1['l_age']=l_age
```


In [63]:

```
AgePlot=df1['age'].sort_values().value_counts()[ :50].plot.bar(figsize=(10,3),x='Value', title="Age"
)
```



In [64]:

```
AgePlot=df1['l_age'].sort_values().value_counts().plot.bar(figsize=(10,3),x='Value', title="Age")
plt.xticks(rotation=45)
plt.show()
```



Workclass

In [65]:

```
RemoveLowCountInstancesCond1=df1['workclass']!="Never-worked"
RemoveLowCountInstancesCond2=df1['workclass']!="Without-pay"
```

In [66]:

```
df1=df1[RemoveLowCountInstancesCond1 & RemoveLowCountInstancesCond2]
```

In [67]:

```
df1.shape
```

Out[67]:

```
(30147, 15)
```

Education

In [68]:

```
df1['education'].value_counts()
```

Out[68]:

```
HS-grad      9831
Some-college 6675
Bachelors    5043
Masters       1627
Assoc-voc    1307
11th         1048
Assoc-acdm   1007
10th          820
7th-8th      556
Prof-school  542
9th          455
12th         377
Doctorate    375
5th-6th      288
1st-4th      151
Preschool    45
Name: education, dtype: int64
```

In [69]:


```
l_education=[]
for elem in df1['education']:
    if elem=="Preschool" or elem=="1st-4th" or elem=="5th-6th":
        l_education.append("Elementary school")
    elif elem=="7th-8th" or elem=="9th" or elem=="10th" or elem=="11th" or elem=="12th":
        l_education.append("HS-Grad_NotCompleted")
    elif elem=="Assoc-voc" or elem=="Assoc-acdm":
        l_education.append("Assoc-acdm_or_Voc")
    else:
        l_education.append(elem)
```


In [70]:

```
df1['l_education']=l_education
```

In [71]:

```
EducationPlot=df1['l_education'].sort_values().value_counts().plot.bar(figsize=(15,5),x='Value', title="Education", grid=True)
```





Marital Status

In [72]:

```
df1['marital-status'].value_counts()
```

Out[72]:

```
Married-civ-spouse      14057
Never-married           9721
Divorced                 4214
Separated                939
Widowed                 826
Married-spouse-absent   369
Married-AF-spouse       21
Name: marital-status, dtype: int64
```

In [73]:

```
MaritalStatusToBeRemoved=df1['marital-status']!="Married-AF-spouse"
```

In [74]:

```
df1=df1[MaritalStatusToBeRemoved]
```

In [75]:


```
df1.shape
```


Out[75]:

```
(30126, 16)
```

In [76]:

```
HoursPerWeekPlot=df1['marital-status'].sort_values().value_counts().plot.bar(figsize=(10,5),x='Value', title="Marital Status")
```





Occupation

In [77]:

```
df1['occupation'].value_counts()
```

Out[77]:

```
Prof-specialty      4035
Craft-repair        4026
Exec-managerial     3991
Adm-clerical        3712
Sales               3582
Other-service       3207
Machine-op-inspct   1965
Transport-moving    1570
Handlers-cleaners   1349
Farming-fishing     982
Tech-support        912
Protective-serv     643
Priv-house-serv     143
Armed-Forces         9
Name: occupation, dtype: int64
```

I decided to remove Armed Forces.

In [78]:

```
df1=df1[df1['occupation']!='Armed-Forces']
```

In [79]:


```
df1['occupation'].value_counts()
```

Out[79]:

```
Prof-specialty      4035
Craft-repair        4026
Exec-managerial     3991
Adm-clerical        3712
Sales               3582
Other-service       3207
Machine-op-inspct   1965
Transport-moving    1570
Handlers-cleaners   1349
Farming-fishing     982
Tech-support        912
Protective-serv     643
Priv-house-serv     143
Name: occupation, dtype: int64
```

In [80]:

```
OccupationPlot=df1['occupation'].sort_values().value_counts().plot.bar(figsize=(15,5),x='Value', title="Occupation")
```



Hours per Week

In [81]:


```
df1['hours-per-week'].sort_values().value_counts()[:20]
```

Out[81]:

```
40    14234
50     2713
45     1753
60     1403
35     1184
20     1052
30      988
55      670
25      570
48      493
38      455
15      346
70      276
65      240
32      238
10      220
24      217
42      212
44      208
36      202
Name: hours-per-week, dtype: int64
```

In [82]:

```
HoursPerWeekPlot=df1['hours-per-week'].sort_values().value_counts()[:20].plot.bar(figsize=(10,3),x
='Value', title="Working Hours per Week")
```



In [83]:

```
l_HoursPerWeek=[]
for elem in df1['hours-per-week']:
    if elem <38:
        l_HoursPerWeek.append("Lower than 38 Hours")
    elif elem >=38 and elem <42:
        l_HoursPerWeek.append(">=38 & <42")


    elif elem >=42 and elem <50:
        l_HoursPerWeek.append(">=42 & <50")
    else:
        l_HoursPerWeek.append("Higher than 50 Hours")
```

In [84]:

```
df1['l_HoursPerWeek']=l_HoursPerWeek
```

In [85]:

```
HoursPerWeekPlotDiscrete=df1['l_HoursPerWeek'].sort_values().value_counts().plot.bar(figsize=(10,3)
,x='Value', title="Working Hours per Week")
plt.xticks(rotation=15)
plt.show()
```



Native Country

In [86]:

```
df1['native-country'].value_counts()[:10]
```

Out[86]:

```
United-States    27460
Mexico           610
Philippines      187
Germany          128
Puerto-Rico     109
Canada           107
India            100
El-Salvador      100
Cuba              92
England           86
Name: native-country, dtype: int64
```

I decided to group in Others all the countries whose observations are lower than 100

In [87]:

```
l_Country=[]
for elem in df1['native-country']:
    if elem == 'United-States':
        l_Country.append(elem)
    else:
```

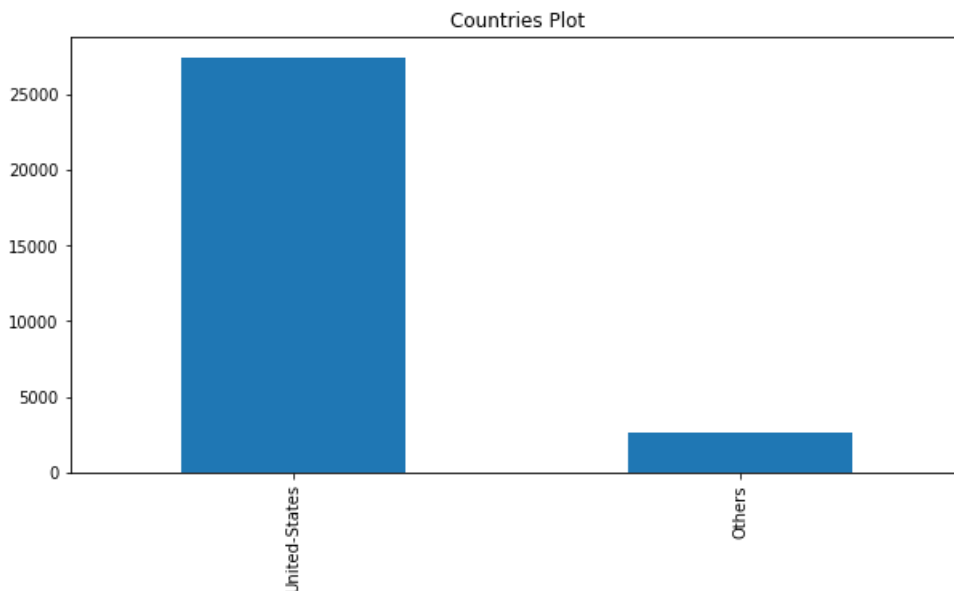
```
l_Country.append('Others')
```

In [88]:

```
df1['l_Country']=l_Country
```

In [89]:

```
Country_Plot=df1['l_Country'].sort_values().value_counts().plot.bar(figsize=(10,5),x='Value', title="Countries Plot")
```



In [90]:

```
df1.head()
```

Out[90]:

	age	workclass	fnlwgt	education	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss	hours-per-week	native-country	Salary	l_Country
0	50	Self-emp-not-inc	83311	Bachelors	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	13	United-States	<=50K	>40K
1	38	Private	215646	HS-grad	Divorced	Handlers-cleaners	Not-in-family	White	Male	0	0	40	United-States	<=50K	>20K
2	53	Private	234721	11th	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	0	0	40	United-States	<=50K	>40K
3	28	Private	338409	Bachelors	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	0	40	Cuba	<=50K	>20K
4	37	Private	284582	Masters	Married-civ-spouse	Exec-managerial	Wife	White	Female	0	0	40	United-States	<=50K	>20K

In [91]:

```
CleanData=['l_age','workclass','l_education','marital-status','occupation','relationship','race','sex','l_Country','Salary']
```

```
In [92]:
```

```
df2=df1[CleanData]
```

```
In [93]:
```

```
df2[:2]
```

```
Out[93]:
```

	l_age	workclass	l_education	marital-status	occupation	relationship	race	sex	l_Country	Salary
0	>40 & <60	Self-emp-not-inc	Bachelors	Married-civ-spouse	Exec-managerial	Husband	White	Male	United-States	<=50K
1	>20 & <40	Private	HS-grad	Divorced	Handlers-cleaners	Not-in-family	White	Male	United-States	<=50K

Latest Check on the content of the attributes

```
In [94]:
```

```
df2['l_age'].value_counts()
```

```
Out[94]:
```

```
>20 & <40      14969
>40 & <60      10315
>80              1760
>60 & <80       1708
<20             1365
Name: l_age, dtype: int64
```

```
In [95]:
```

```
df2['workclass'].value_counts()
```

```
Out[95]:
```

```
Private      22271
Self-emp-not-inc  2497
Local-gov    2067
State-gov    1277
Self-emp-inc  1074
Federal-gov   931
Name: workclass, dtype: int64
```

```
In [96]:
```

```
df2['l_education'].value_counts()
```

```
Out[96]:
```

```
HS-grad      9815
Some-college 6670
Bachelors    5038
HS-Grad_NotCompleted 3255
Assoc-acdm_or_Voc 2312
Masters      1626
Prof-school   542
Elementary school 484
Doctorate     375
Name: l_education, dtype: int64
```

```
In [97]:
```

```
df2['marital-status'].value_counts()
```

Out[97]:

```
Married-civ-spouse      14054
Never-married           9715
Divorced                4214
Separated               939
Widowed                 826
Married-spouse-absent   369
Name: marital-status, dtype: int64
```

In [98]:

```
df2['occupation'].value_counts()
```

Out[98]:

```
Prof-specialty          4035
Craft-repair            4026
Exec-managerial         3991
Adm-clerical            3712
Sales                   3582
Other-service           3207
Machine-op-inspct       1965
Transport-moving        1570
Handlers-cleaners       1349
Farming-fishing          982
Tech-support            912
Protective-serv         643
Priv-house-serv         143
Name: occupation, dtype: int64
```

In [99]:

```
df2['relationship'].value_counts()
```

Out[99]:

```
Husband                12448
Not-in-family          7721
Own-child              4459
Unmarried              3210
Wife                   1393
Other-relative         886
Name: relationship, dtype: int64
```

In [100]:

```
df2['race'].value_counts()
```

Out[100]:

```
White                25893
Black                2814
Asian-Pac-Islander   894
Amer-Indian-Eskimo   285
Other                231
Name: race, dtype: int64
```

In [101]:

```
df2['sex'].value_counts()
```

Out[101]:

```
Male      20352
Female     9765
Name: sex, dtype: int64
```

In [102]:

```
df2['l_Country'].value_counts()
```

Out[102]:

```
United-States    27460
Others           2657
Name: l_Country, dtype: int64
```

In [103]:

```
df2['Salary'].value_counts()
```

Out[103]:

```
<=50K    22620
>50K      7497
Name: Salary, dtype: int64
```

In []:

It seems to be good. Now we can start with the classification activities

In [104]:

```
df2[-2:]
```

Out[104]:

	l_age	workclass	l_education	marital-status	occupation	relationship	race	sex	l_Country	Salary
32558	>20 & <40	Private	HS-grad	Never-married	Adm-clerical	Own-child	White	Male	United-States	<=50K
32559	>40 & <60	Self-emp-inc	HS-grad	Married-civ-spouse	Exec-managerial	Wife	White	Female	United-States	>50K

Now, I need to encode correctly my dataset

In [105]:

```
le = preprocessing.LabelEncoder()
df3 = df2.apply(le.fit_transform)
```

In [106]:

```
df3[:20]
```

Out[106]:

	l_age	workclass	l_education	marital-status	occupation	relationship	race	sex	l_Country	Salary
0	2	4	1	1	2	0	4	1	1	0
1	1	2	5	0	4	1	4	1	1	0
2	2	2	4	1	4	0	2	1	1	0
3	1	2	1	1	8	5	2	0	0	0
4	1	2	6	1	2	5	4	0	1	0
5	2	2	4	2	6	1	2	0	0	0
6	2	4	5	1	2	0	4	1	1	1
7	1	2	6	3	8	1	4	0	1	1
8	2	2	1	1	2	0	4	1	1	1

	l_age	workclass	l_education	marital-status	occupation	relationship	race	sex	l_Country	Salary
9	1	5	1	1	8	0	1	1	0	1
10	1	2	1	3	0	3	4	0	1	0
11	1	2	0	3	10	1	2	1	1	0
12	1	2	4	1	12	0	0	1	0	0
14	1	4	5	3	3	3	4	1	1	0
15	1	2	5	3	5	4	4	1	1	0
16	1	2	4	1	10	0	4	1	1	0
17	2	4	6	0	2	4	4	0	1	1
18	4	2	2	1	8	0	4	1	1	1
19	2	2	5	4	6	4	2	0	1	0
20										

- I create my target Variable

In [107]:

```
y=df3[['Salary']].copy()
y
```

Out[107]:

	Salary
0	0
1	0
2	0
3	0
4	0
...	...
32555	0
32556	1
32557	0
32558	0
32559	1

30117 rows × 1 columns

In [108]:

```
X=df3.iloc[:,0:9]
```

In [109]:

```
X.columns
```

Out[109]:

```
Index(['l_age', 'workclass', 'l_education', 'marital-status', 'occupation',
      'relationship', 'race', 'sex', 'l_Country'],
      dtype='object')
```

Training Phase

In [110]:

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state=324)
```

In [111]:

```
#type(X_train)
#type(X_test)
#type(y_train)
#type(y_test)
#X_train.head()
#y_train.describe()
#X_test.describe()
```

In [112]:

```
IncomeClassifier = DecisionTreeClassifier(random_state=0,max_leaf_nodes=20)
IncomeClassifier.fit(X_train, y_train)
```

Out[112]:

```
DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=None,
                        max_features=None, max_leaf_nodes=20,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, presort=False,
                        random_state=0, splitter='best')
```

In [113]:

```
type(IncomeClassifier)
```

Out[113]:

```
sklearn.tree.tree.DecisionTreeClassifier
```


In [114]:

```
print(IncomeClassifier.feature_importances_)
```

```
[0.03562728 0.00585681 0.28184603 0.          0.05089583 0.62577404
 0.          0.          0.          ]
```

In [115]:

```
pyplot.bar(range(len(IncomeClassifier.feature_importances_)),
IncomeClassifier.feature_importances_)
pyplot.show()
```



In [116]:

```
names=X.columns
```


In [117]:

```
pyplot.bar(range(len(IncomeClassifier.feature_importances_)),
IncomeClassifier.feature_importances_)
```

```

pyplot.xticks(range(X.shape[1]), names, rotation=-60)
pyplot.title('Most Informative Variables')
pyplot.grid(True)
pyplot.show()

```



I print the decision tree

In [118]:

```

dot_data = tree.export_graphviz(IncomeClassifier, out_file=None, filled = True, feature_names=names, class_names=df2['Salary'].unique())
graph = graphviz.Source(dot_data)
pydot_graph = pydotplus.graph_from_dot_data(dot_data)
pydot_graph.write_png('original_tree.png')
pydot_graph.set_size('10,10!')
pydot_graph.write_png('resized_tree.png')

```

Out[118]:

True

In [119]:

graph

Out[119]:



In []:

Prediction on Test Set

In [120]:

```

predictions = IncomeClassifier.predict(X_test)

```

In [121]:

```

predictions[-20:]

```

Out[121]:

```

array([1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0])

```

In [122]:

```
y_test['Salary'][-20:]
```

Out[122]:

```
3322      0
31131     0
19679     0
8480      0
9506      0
23265     1
17814     0
7524      0
24188     0
18555     0
994       1
2116      0
22198     0
6641      0
27912     0
5270      0
27457     1
28073     0
6593      0
18715     0
Name: Salary, dtype: int32
```

Measure of accuracy

In [123]:

```
accuracy_score(y_true = y_test,
               y_pred = predictions)
```

Out[123]:

```
0.820303853506389
```

In []: